

# Reinforcement Learning David Silver - Lecture 7 Notes: Policy Gradient

Name: Eli Andrew

## • Policy-based RL

- Previously, we were parameterizing  $V$  and  $Q$  and generating a policy directly from them (using  $\epsilon$ -greedy for example)
- Now, we directly parameterize the policy

$$\pi_{\theta}(s, a) = P[a|s, \theta]$$

- The motivation here is, once again, to be able to scale efficiently to many states
- Advantages:
  - \* Better convergence properties
  - \* (best reason) Effective in high-dimensional or continuous action spaces
  - \* Can learn stochastic policies
- Disadvantages:
  - \* Typically converge to a local rather than global optimum
  - \* Evaluating a policy is typically inefficient and high variance
- Why would you want a stochastic policy?
  - \* Games like Rock-paper-scissors
  - \* State aliasing problems where the agent can't differentiate between certain states
  - \* The state aliasing problem can occur because of a partially observed environment (which is equivalent to not having the correct features to represent the environment)
  - \* If we have state aliasing with a deterministic policy, then all aliased states (different states that appear the same) will have to have the same action

## • Policy Objective Functions

- Goal: given policy  $\pi_{\theta}(s, a)$  with parameters  $\theta$ , find best  $\theta$
- How to measure quality of policy  $\pi_{\theta}(s, a)$ 
  - \* Episodic environments: use **start value**

$$J_1(\theta) = V^{\pi_{\theta}}(s_1) = E_{\pi_{\theta}}[v_1]$$

- \* In continuing environments: use **average reward per time-step**

$$J_{avR}(\theta) = \sum_s d^{\pi_{\theta}}(s) \sum_a \pi_{\theta}(s, a) R_s^a$$

- \* Or **average value**

$$J_{avV}(\theta) = \sum_s d^{\pi_\theta}(s) V^{\pi_\theta}(s)$$

- \* Where  $d^{\pi_\theta}(s)$  is stationary distribution of Markov chain for  $\pi_\theta$

## • Policy Optimization

- Find  $\theta$  that minimises  $J(\theta)$
- Some approaches don't use gradient
  - \* Hill climbing
  - \* Simplex / amoeba / Nelder Mead
  - \* Genetic algorithms
- Greater efficiency often possible using gradient
  - \* Gradient descent
  - \* Conjugate gradient
  - \* Quasi-newton

## • Policy Gradient

- Let  $J(\theta)$  be any policy objective function
- Policy gradient algorithms search for *local* maximum in  $J(\theta)$  by ascending the gradient of the policy, w.r.t parameters  $\theta$

$$\Delta\theta = \alpha \nabla_\theta J(\theta)$$

- Where  $\nabla_\theta J(\theta)$  is the **policy gradient** and  $\alpha$  is step size parameter

## • Score Function

- Assume policy  $\pi_\theta$  is differentiable everywhere we can take actions
- Assume we know gradient  $\nabla_\theta \pi_\theta(s, a)$
- Likelihood ratios exploit the following identity:

$$\begin{aligned} \nabla_\theta \pi_\theta(s, a) &= \pi_\theta(s, a) \frac{\nabla_\theta \pi_\theta(s, a)}{\pi_\theta(s, a)} \\ &= \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a) \end{aligned}$$

- Score Function is:  $\nabla_\theta \log \pi_\theta(s, a)$
- Using this form allows us to take expectations much easier

## • Softmax Policy

- Weight actions using linear combination of features  $\phi(s, a)^\top \theta$
- Probability of action is proportional to exponentiated weight:

$$\pi_\theta(s, a) \propto \exp(\phi(s, a)^\top \theta)$$

- Score function is then:

$$\nabla_\theta \log \pi_\theta(s, a) = \phi(s, a) - E_{\pi_\theta}[\phi(s, \cdot)]$$

- In other words the score function is the difference between the feature of the action we actually took and the average feature: this says how much more of this function do I have than usual
- Since adjustments are then made based on this difference, it means that if I feature contributed more than usual then it will be updated more in the direction of the reward we received

### • Gaussian Policy

- Most common policy for continuous action spaces
- Mean is linear combination of state features

$$\mu(s) = \phi(s)^\top \theta$$

- Variance may be fixed  $\sigma^2$  or can be parameterized
- Policy is Gaussian:

$$a \sim \mathcal{N}(\mu(s), \sigma^2)$$

- Score function is:

$$\nabla_\theta \log \pi_\theta(s, a) = \frac{(a - \mu(s))\phi(s)}{\sigma^2}$$

### • One step MDP

- Use example of a simple one-step MDP where you start in state  $s \sim d(s)$  and terminate after one step with reward  $r = R_{s,a}$
- Likelihood ratios help with calculating Policy gradient:

$$\begin{aligned} J(\theta) &= E[r] \\ &= \sum_{s \in S} d(s) \sum_{a \in A} \pi_\theta(s, a) R_{s,a} \\ \nabla_\theta J(\theta) &= \sum_{s \in S} d(s) \sum_{a \in A} \pi_\theta(s, a) \nabla_\theta \log(\pi_\theta(s, a)) R_{s,a} \\ &= E[\nabla_\theta \log(\pi_\theta(s, a)) r] \end{aligned}$$

- **Policy Gradient Theorem**

- Generalizes likelihood ratio to multi-step MDPs
- Replaces instantaneous reward  $r$  with long term value  $Q^\pi(s, a)$
- Applies to start state objective, average reward, and average value objective
- Theorem: for any differentiable policy  $\pi_\theta(s, a)$  for any of the applicable objective functions (stated above) the policy gradient is:

$$\nabla_\theta J(\theta) = E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) Q^{\pi_\theta}(s, a)]$$

- **Monte Carlo Policy Gradient**

- Use  $v_t$  as an unbiased sample of  $Q^{\pi_\theta}(s_t, a_t)$
- For each episode: select  $s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T \sim \pi_\theta$
- For each time step of the episode make update:  $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$
- Where  $v_t$  is cumulated rewards from that time step onwards

- **Reducing variance with a Critic**

- Monte-Carlo policy gradient has high variance
- Use a **critic** to estimate the state-action value function:

$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$

- Actor-critic algorithms maintain two sets of parameters
  - \* Critic: Updates action-value function parameters  $w$
  - \* Actor: Updates policy parameters  $\theta$ , in direction suggested by critic
- Actor-critic algorithms follow an *approximate* policy gradient

$$\begin{aligned} \nabla_\theta J(\theta) &\approx E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)] \\ \Delta\theta &= \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a) \end{aligned}$$

- **Action-Value Actor-Critic**

- Using linear value function approx.  $Q_w(s, a) = \phi(s, a)^\top w$ 
  - \* Critic: Updates  $w$  by linear TD(0)
  - \* Actor: Updates  $\theta$  by policy gradient
  - \* Initialize  $s, \theta$
  - \* Sample  $a \sim \pi_\theta$

- \* For each step:

Sample reward  $r = R_s^a$ ; Sample transition  $s' \sim P_s^a$

Sample action  $a' \sim \pi_\theta(s', a')$

$\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

$\theta = \theta + \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$

$w \leftarrow w + \beta \delta \phi(s, a)$

$a \leftarrow a', s \leftarrow s'$

- In summary:

- \* Actor picks the actions using some policy
- \* Critic evaluates and says whether the actions are good or bad
- \* Actor moves its policy in the direction suggested by the critic

### • Reducing Variance Using a Baseline

- Subtract a baseline function  $B(s)$  from the policy gradient
- This reduces variance without changing expectation:

$$\begin{aligned} E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) B(s)] &= \sum_{s \in S} d^{\pi_\theta}(s) \sum_a \nabla_\theta \pi_\theta(s, a) B(s) \\ &= \sum_{s \in S} d^{\pi_\theta}(s) \nabla_\theta \sum_{a \in A} \pi_\theta(s, a) B(s) \\ &= 0 \end{aligned}$$

- State value function is good baseline:  $B(s) = V^{\pi_\theta}(s)$
- Re-write policy gradient using advantage function:  $A^{\pi_\theta}(s, a)$ :

$$\begin{aligned} A^{\pi_\theta}(s, a) &= Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s) \\ \nabla_\theta J(\theta) &= E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) A^{\pi_\theta}(s, a)] \end{aligned}$$

- Advantage function tells us how much better it is than usual to take action  $a$
- $\nabla_\theta \log \pi_\theta(s, a)$  tells us how to adjust  $\theta$  to achieve that action  $a$  with our policy  $\pi_\theta$

### • Estimating the Advantage Function

- TD-error is a sample of the advantage function
- For true value function  $V^{\pi_\theta}(s)$ , the TD-error  $\delta^{\pi_\theta}$

$$\delta^{\pi_\theta} = r + \gamma V^{\pi_\theta}(s') - V^{\pi_\theta}(s)$$

- is an unbiased estimate of the advantage function:

$$\begin{aligned} E[\delta^{\pi_\theta} | s, a] &= E_{\pi_\theta}[r + \gamma V^{\pi_\theta}(s') | s, a] - V^{\pi_\theta}(s) \\ &= Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s) = A^{\pi_\theta}(s, a) \end{aligned}$$

- So, we can use TD error to compute policy gradient:

$$\nabla_\theta J(\theta) = E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \delta^{\pi_\theta}]$$

- And in practice we can just estimate the TD error:

$$\delta_v = r + \gamma V_v(s') - V_v(s)$$

- This leads to only one set of parameters:  $v$

### • Natural Policy Gradient

- Policies we've considered have all been stochastic
- We've estimated our policy gradients by sampling our own noise
- Taking expectation of gradient of our own noise can run into issues, especially as Gaussian gets narrower as your policy improves (noise hurts you more the better the policy gets)
- Instead, start with a deterministic policy
- With the deterministic case where our noise is narrowed down to 0 we get an update:

$$\nabla_\theta^{\text{nat}} \pi_\theta(s, a) = G_\theta^{-1} \nabla_\theta \pi_\theta(s, a)$$

- Where  $G_\theta$  is Fischer Information matrix:  $E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)^\top]$
- In practice, this performs much better than stochastic gradient policies especially in cases with continuous actions

### • Natural Actor-Critic

- Using compatible function approximation

$$\nabla_w A_w(s, a) = \nabla_\theta \log \pi_\theta(s, a)$$

- Natural policy gradient becomes:

$$\begin{aligned} \nabla_\theta J(\theta) &= E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) A^{\pi_\theta}(s, a)] \\ &= E_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)^\top w] \\ &= G_\theta w \nabla_\theta^{\text{nat}} J(\theta) = w \end{aligned}$$

- **Summary of Policy Gradient Algorithms**

- The **policy gradient** has many equivalent forms

- REINFORCE:  $v_t$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) v_t]$$

- Q Actor-Critic:  $Q^w(s, a)$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) Q^w(s, a)]$$

- Advantage Actor-Critic:  $A^w(s, a)$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) A^w(s, a)]$$

- TD Actor-Critic:  $\delta$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) \delta]$$

- TD( $\lambda$ ) Actor-Critic:  $\delta e$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) \delta e]$$

- Natural Actor-Critic:

$$G_{\theta}^{-1} \nabla_{\theta} J(\theta) = w$$

- Each form leads a stochastic gradient ascent algorithm

- Critic uses **policy evaluation** (e.g. MC, TD, or other algorithms from before) to estimate  $Q^{\pi}(s, a)$ ,  $A^{\pi}(s, a)$ , or  $V^{\pi}(s)$