Reinforcement Learning
Name: Eli Andrew

Exercise 1.1: Self-Play

(a) **Question:** Suppose, intead of playing against a random opponent, the reinforcement learning algorithm described above played against itself, with both sides learning. What do you think would happen in this case? Would it learn a different policy for selecting moves?